

# Reward prediction error neurons implement an efficient code for reward

Authors: Heiko H. Schütt, Dongjae Kim, Wei Ji Ma

Presented by Snow

# Efficient Coding

sensory neuroscience

neuronal codes are optimized for efficiency, that is, to convey as much **information** as possible with **given budgets** for the number of neurons and the number of action potentials

# Efficient Coding

sensory neuroscience

neuronal codes are optimized for efficiency, that is, to convey as much **information** as possible with **given budgets** for the number of neurons and the number of action potentials



reward

Reward Prediction Error Neurons?  
(RPEN)

# Overview

1. Derive the efficient code
2. Efficient code satisfies properties of real neurons
3. How do neurons learn the tuning? Distributional RL

# Overview

## 1. Derive the efficient code

### **Efficient Sensory Encoding and Bayesian Inference with Heterogeneous Neural Populations**

**Deep Ganguli and Eero P. Simoncelli**

Howard Hughes Medical Institute, Center for Neural Science, and Courant Institute of Mathematical Sciences, New York University, New York, NY 10003

Deep Ganguli: [dganguli@cns.nyu.edu](mailto:dganguli@cns.nyu.edu); Eero P. Simoncelli: [eero.simoncelli@nyu.edu](mailto:eero.simoncelli@nyu.edu)


with some modifications

# Overview

1. Derive the efficient code


Main idea: Optimize Fisher Information (with certain constraints)

# Efficient Code computation



RPEN's sigmoid tuning curve

## What are the elements of a population coding?



RPEN's sigmoid tuning curve


(given N neurons)

- Tuning curves(response functions)  
(assume they are the same w.r.t. given midpoint= $\mu$ )
- Density of neurons with midpoint= $\mu$
- Regularization: Mapping from reward to unit interval

$$h_\mu(R) = h_0[F(R) - F(\mu)]g(\mu)$$

## Efficient Code computation

$$h_\mu(R) = h_0[F(R) - F(\mu)]g(\mu)$$




RPEN's sigmoid tuning curve

- Response function, midpoint =  $\mu$
- Reward mapping  $F : R \rightarrow [0, 1]$
- A prototypical sigmoidal function  $h_0$
- Neural Gain  $g$  as a function of  $R$


## Efficient Code computation

Problem of efficient code:




What is the most efficient population to encode the reward stimuli?

Problem of efficient code:




What is the most efficient population to encode the reward stimuli?

## Problem of efficient code:



What is the most efficient population to encode the reward stimuli?

## Problem of efficient code:



Fisher Information




maximize  
under  
constraints:

upper bound to the expected  
number of spikes

What is the most efficient  
population to encode the reward  
stimuli?

## Problem of efficient code:



Fisher Information



maximize  
under  
constraints:

upper bound to the expected  
number of spikes

$$f(R) = p(R)$$

$$d(R) \propto \frac{p(R)}{[1 - P(R)]^{1-\alpha}}$$

$$g(R) \propto \frac{1}{[1 - P(R)]^\alpha}$$

$\alpha$  is a free parameter,  $\alpha \in [0, 1]$

What is the most efficient  
population to encode the reward  
stimuli?


Zoom back into a single neuron:  $h_\mu(R) = h_0[F(R) - F(\mu)]g(\mu)$

$$f(R) = p(R)$$

$$d(R) \propto \frac{p(R)}{[1 - P(R)]^{1-\alpha}}$$

$$g(R) \propto \frac{1}{[1 - P(R)]^\alpha}$$


$\alpha$  is a free parameter,  $\alpha \in [0, 1]$



these results are **independent** of  
the exact shape of the sigmoid  
tuning curves( $h_0$ )

## Efficient Code computation

Zoom back into a single neuron:  $h_\mu(R) = h_0[F(R) - F(\mu)]g(\mu)$




slope  $k$ ; spontaneous firing rate  $r^*$

we fit these two parameters:  
minimize MSE of thresholds  
(between efficient code and measured neurons)

c.d.f. of beta distributions as the  
sigmoid response functions

# Cumulative Distribution Function of Beta Distribution



Efficient Code computation


Zoom back into a single neuron:

c.d.f. of beta distributions as the sigmoid response functions

beta distribution is parameterized  
by  $a, b$

for a neuron with slope  $k$ , midpoint  $p$

set  $a = kp$ ,  $b = k(1-p)$



good boundary behavior

# What's new?

## **Efficient Sensory Encoding and Bayesian Inference with Heterogeneous Neural Populations**

**Deep Ganguli and Eero P. Simoncelli**

Howard Hughes Medical Institute, Center for Neural Science, and Courant Institute of Mathematical Sciences, New York University, New York, NY 10003

Deep Ganguli: [dganguli@cns.nyu.edu](mailto:dganguli@cns.nyu.edu); Eero P. Simoncelli: [eero.simoncelli@nyu.edu](mailto:eero.simoncelli@nyu.edu)

$$f(R) = p(R)$$

$$d(R) \propto \frac{p(R)}{[1 - P(R)]^{1-\alpha}} \quad \text{a special case: } \alpha = 1$$


$$g(R) \propto \frac{1}{[1 - P(R)]^\alpha}$$

# Overview

1. Derive the efficient code
2. Efficient code satisfies properties of real neurons
3. How do neurons learn the tuning? Distributional RL

# Application 1: variable-reward task

reward condition:



vertical bars: seven values of reward  
with probability

smooth curve: fitted moment-  
matched log-normal distribution

Application 1: variable-reward task

## Data & Model fitting

Application 1: variable-reward task

## Data & Model fitting

Step 1: Fitting Logistic sigmoid functions to measured neurons

$$\sigma(R) = \frac{a_1}{1 - e^{-a_2(R-a_3)}}$$

maximum-likelihood estimation based on Poisson noise

$$\hat{\mathbf{a}} = \operatorname{argmin}_{\mathbf{a}} \sum_{j=1}^T (r_j \log(\sigma(R_j; \mathbf{a})) - \sigma(R_j; \mathbf{a}))$$

where  $\mathbf{a} = (a_1, a_2, a_3)$

Application 1: variable-reward task

## Data & Model fitting

Step 2:  $\alpha$

$$f(R) = p(R)$$

$$d(R) \propto \frac{p(R)}{[1 - P(R)]^{1-\alpha}}$$

$$g(R) \propto \frac{1}{[1 - P(R)]^\alpha}$$

$\alpha$  is a free parameter,  $\alpha \in [0, 1]$

Application 1: variable-reward task

## Data & Model fitting

Step 2:  $\alpha$

maximize the probability of the observed midpoints

Application 1: variable-reward task

## Data & Model fitting

Step 2:  $\alpha$

maximize the probability of the observed midpoints


$$\alpha = 0.673$$

## Data & Model fitting

Step 3: slope  $k$ ; spontaneous firing rate  $r^*$


minimize MSE of thresholds  
(between efficient code and  
measured neurons)

$$\operatorname{argmin}_{r^*, k} E \left[ \theta_i - h_i^{-1}(r^*, k) \right]^2$$



**c**  
Application 1: variable-reward task


Response (Hz)



39 dopaminergic RPENs  
measured by Eshel et al.  
(preprocessed)

**d**

Response (Hz)



Efficient code

## Application 1: variable-reward task

### ***Predictions:***



1. RPEN midpoints cover the reward distribution
2. RPEN gain increases with threshold
3. RPEN tuning curve asymmetry flips with increasing threshold
4. RPEN slope decreases with threshold

My question: Why \*these\* predictions?

## Application 1: variable-reward task

### Predictions

Prediction 1: RPEN midpoints cover the reward distribution




difference not significant

## Application 1: variable-reward task

### Predictions

Prediction 2: RPEN gain increases with threshold



with same expected number of spikes:

higher threshold  $\rightarrow$  respond to fewer rewards

afford a higher gain

Blue: Efficient code


Black: Measured neurons

this empirical finding has not been reported before

## Application 1: variable-reward task

### Predictions


Prediction 3: RPEN tuning curve asymmetry flips with increasing threshold



## Application 1: variable-reward task

### Predictions

Prediction 3: RPEN tuning curve asymmetry flips with increasing threshold



Degree of asymmetry:


$$\frac{\beta^+}{\beta^+ + \beta^-} \rightarrow \begin{array}{l} \text{slope above threshold} \\ \text{slope below threshold} \end{array}$$

Fit two linear response functions  
(following Dabney's work)

## Application 1: variable-reward task

### Predictions

Prediction 3: RPEN tuning curve asymmetry flips with increasing threshold




strong positive relationship

## Application 1: variable-reward task

### Predictions

Prediction 3: RPEN tuning curve asymmetry flips with increasing threshold


What is the meaning of asymmetry here?



## Application 1: variable-reward task

### Predictions

**Prediction 4: RPEN slope decreases with threshold**




response functions are **shallower** in regions with **lower** reward density  
( $\Rightarrow$ cover broader reward ranges with fewer neurons)


this empirical finding has not been reported before

# Application 2: variable-distribution task

reward condition: Rothenhoefer et al.



‘normal’ distribution  
(with the same mean)




‘uniform’ distribution

## Application 2: variable-distribution task

example neurons

c




main finding of Rothenhoefer et al. :

steeper response functions for  
'normal' than 'uniform'

(dopamine responses are amplified  
for rare rewards)

## Application 2: variable-distribution task

all neurons



steeper response functions for  
'normal' than 'uniform'

# Overview

1. Derive the efficient code
2. Efficient code satisfies properties of real neurons
3. How do neurons learn the tuning? Distributional RL

# Reinforcement Learning

$$Q(x, a) = \mathbb{E}R(x, a) + \gamma\mathbb{E}Q(X', A')$$

# Distributional Reinforcement Learning

$$Z(x, a) \stackrel{D}{=} R(x, a) + \gamma Z(X', A')$$



random variable instead of an expected value

---


## A Distributional Perspective on Reinforcement Learning

---

Marc G. Bellemare <sup>\* 1</sup> Will Dabney <sup>\* 1</sup> Rémi Munos <sup>1</sup>

# Distributional Reinforcement Learning

$$Z(x, a) \stackrel{D}{=} R(x, a) + \gamma Z(X', A')$$



# Distributional Reinforcement Learning

## The Wasserstein Metric

calculate the distance between two distributions

$$d_p(F, G) := \inf_{U, V} \|U - V\|_p$$


# Distributional Reinforcement Learning

How to implement in tasks?

# Distributional Reinforcement Learning

How to implement in tasks?


Approximate the distribution(continuous)  
with a **discrete** set of parametrized value  
functions



“highly expressive and  
computationally friendly”

# Distributional Reinforcement Learning

$$Z(x, a) \stackrel{D}{=} R(x, a) + \gamma Z(X', A')$$



DRL in dopamine neurons

Article

# A distributional code for value in dopamine-based reinforcement learning

---

<https://doi.org/10.1038/s41586-019-1924-6>

---



Received: 3 January 2019

---

Accepted: 19 November 2019

Will Dabney<sup>1,5\*</sup>, Zeb Kurth-Nelson<sup>1,2,5</sup>, Naoshige Uchida<sup>3</sup>, Clara Kwon Starkweather<sup>3</sup>, Demis Hassabis<sup>1</sup>, Rémi Munos<sup>1</sup> & Matthew Botvinick<sup>1,4,5</sup>

# DRL in dopamine neurons



For each value function  $V_i$ , prediction error:

$$\delta_i = r + \gamma V_j(x') - V_i(x)$$

Asymmetry update rule:

$$V_i(x) \leftarrow V_i(x) + \alpha_i^+ \delta_i \quad \text{for } \delta_i > 0$$

$$V_i(x) \leftarrow V_i(x) + \alpha_i^- \delta_i \quad \text{for } \delta_i < 0$$

## DRL in dopamine neurons

$$V_i(x) \leftarrow V_i(x) + \alpha_i^+ f(\delta_i) \quad \text{for } \delta_i > 0$$

$$V_i(x) \leftarrow V_i(x) + \alpha_i^- f(\delta_i) \quad \text{for } \delta_i < 0$$

where  $f$  is a function that transforms the prediction error

## DRL in dopamine neurons

$$V_i(x) \leftarrow V_i(x) + \alpha_i^+ f(\delta_i) \quad \text{for } \delta_i > 0$$

$$V_i(x) \leftarrow V_i(x) + \alpha_i^- f(\delta_i) \quad \text{for } \delta_i < 0$$

take  $f(x) = sign(x)$   $\alpha_i^+ = 3, \alpha_i^- = 1$

## DRL in dopamine neurons


$$V_i(x) \leftarrow V_i(x) + 3 \quad \text{for } \delta_i > 0$$

$$V_i(x) \leftarrow V_i(x) - 1 \quad \text{for } \delta_i < 0$$

## DRL in dopamine neurons

$$V_i(x) \leftarrow V_i(x) + 3 \quad \text{for } \delta_i > 0$$

$$V_i(x) \leftarrow V_i(x) - 1 \quad \text{for } \delta_i < 0$$





## DRL in dopamine neurons

$$V_i(x) \leftarrow V_i(x) + \alpha_i^+ \quad \text{for } \delta_i > 0$$

$$V_i(x) \leftarrow V_i(x) - \alpha_i^- \quad \text{for } \delta_i < 0$$

$$\tau_i = \frac{\alpha_i^+}{\alpha_i^+ + \alpha_i^-}$$




## Application 1: variable-reward task

### Predictions

Prediction 3: RPEN tuning curve asymmetry flips with increasing threshold

What is the meaning of asymmetry here?



# Learning rules for the efficient code

□ Quantile learning

Application 1: variable-reward task

□ Slope learning

*Predictions:*


□ Gain

1. RPEN midpoints cover the reward distribution
2. RPEN gain increases with threshold
3. RPEN tuning curve asymmetry flips with increasing threshold
4. RPEN slope decreases with threshold

## Learning rules for the efficient code

### Quantile learning

**a**



If reward > midpoint, shift the curve to the right,  
step size = **b**


If reward < midpoint, shift the curve to the left,  
step size = **a**

converge to the **p**th quantile of the reward  
distribution

where **ap = b(1 - p)**

## Learning rules for the efficient code

### Expectile learning




If reward > midpoint, shift the curve to the right,  
step size = **b** \* prediction error

If reward < midpoint, shift the curve to the left,  
step size = **a** \* prediction error

## Learning rules for the efficient code

### Slope learning

**b**




In the efficient code, the RPEN slope should be proportional to the local probability density of rewards

Learning rules for the efficient code

## Gain Method 1


### Gain learning



## Learning rules for the efficient code


### Gain Method 2

Fixed gain per neuron, determined by quantile



## Learning rules for the efficient code

### Comparison



# Efficient Code computation Details

## Efficient Code computation

We start by defining a ‘standard’ population of RPENs

## Efficient Code computation


# Discussion

Task design: No action required by the animal



the reward received and the change in value of the current state are confounded

# Discussion



$$h_\mu(R) = h_0[F(R) - F(\mu)]g(\mu)$$

- Response function, midpoint =  $\mu$
- Reward mapping  $F : R \rightarrow [0, 1]$
- A prototypical sigmoidal function  $h_0$
- Neural Gain  $g$  as a function of  $R$

We assumed that every neuron with a midpoint =  $\mu$  is the same.

# Discussion

In classical reinforcement learning, we set the same learning rate for negative PE and positive PE.

What about setting two different ones respectively? Would that help in better fitting behavioral data, or extract some latent variables for further analysis?